Evaluator Answer Key

Keep this separate from the PDF uploaded to the app. Fixture NS-2026-01.

Score each question out of 2: 1 point for a correct, complete answer and 1 point for a valid source-page/section citation. Maximum: 30 points. For Q15, accept equivalent wording that includes all three required tradeoffs. Numerical equivalence is acceptable unless a precision is specified.

An unsupported factual answer to Q13 or Q14 is a hallucination flag regardless of total score. Track answer accuracy (15 points) separately from citation accuracy (15 points). This small synthetic test is diagnostic; it does not establish general model accuracy.

Q01 | Who owns the pilot, and what dates did it cover?

Maya Chen; September 1-30, 2026.

Evidence: source PDF page 1, section 1.

Q02 | What percentage of requests were extraction tasks?

250 / 1,000 x 100 = 25%.

Evidence: source PDF page 1, section 1.

Q03 | What were the candidate request-success rate and all-attempt quality-pass rate?

Success: 980 / 1,000 = 98%. Quality: 940 / 1,000 = 94%.

Evidence: source PDF page 1, section 2.

Q04 | Does the candidate meet the release quality threshold?

No. 94% is below 94.5% by 0.5 percentage points.

Evidence: source PDF page 1, section 2; 2, section 5.

Q05 | What were the modeled cost saving per 1,000 requests and percentage reduction?

$12.50 - $8.00 = $4.50; $4.50 / $12.50 x 100 = 36%.

Evidence: source PDF page 1, section 2; 2, section 4.

Q06 | How much did running both configurations actually cost?

$12.50 + $8.00 = $20.50 evaluation expenditure.

Evidence: source PDF page 1, section 2.

Q07 | What is the candidate cost per quality-passing response, rounded to six decimal places?

$8.00 / 940 = $0.008511 per quality-passing response.

Evidence: source PDF page 1, section 2.

Q08 | Did candidate median latency and 95th-percentile latency both improve?

No. Median decreased 180 ms (820 - 640); p95 increased 200 ms (1,900 - 1,700).

Evidence: source PDF page 1, section 2.

Q09 | For NS-041, report uncached input, cached input, output, reasoning and total tokens.

9,000 uncached; 3,000 cached; 2,000 output including 500 reasoning; 14,000 total.

Evidence: source PDF page 2, section 3.

Q10 | Calculate the full NS-041 API cost and each component.

Uncached $0.001800; cache read $0.000150; output $0.001600; total $0.003550. No extra reasoning charge.

Evidence: source PDF page 2, section 3.

Q11 | Should the candidate pilot cost be reported as $7.60 or $8.00, and why?

$8.00. Revision 2 adds $0.40 failed-attempt charges and supersedes revision 1.

Evidence: source PDF page 2, section 4.

Q12 | Under the stated 100,000-request projection, what are baseline cost, candidate cost and savings?

$1,250 baseline; $800 candidate; $450 savings. Multiply pilot costs by 100. Hypothetical, not observed.

Evidence: source PDF page 1, section 2; 2, section 5.

Q13 | How many kWh of measured electricity did the candidate save?

Not provided in the document. Do not infer electricity from API cost.

Evidence: source PDF page 2, sections 4-5.

Q14 | What GPU type and provider region were used?

Not provided in the document for either field.

Evidence: source PDF page 2, section 4.

Q15 | Summarize whether the candidate should be released under the stated policy in two sentences.

Must mention lower cost but failure of the quality threshold; also distinguish improved median from worse p95. Do not recommend release as policy-compliant.

Evidence: source PDF page 1, section 2; 2, section 5.

Testing procedure

Upload only the source PDF. Submit Q01-Q15 individually or as one batch using the instructions on page 3. Save model ID, full answer, citations, tokens, latency and cost. Test each model against the same PDF and questions. Compare against this key manually; do not treat successful HTTP requests as correct answers.