Northstar Support Pilot

Accuracy test source | Fixture NS-2026-01 | Revision 2

This fictional report is designed to test document extraction, grounded answers and simple calculations. All company names, rates and measurements are synthetic. These rates are not current provider pricing.

1. Scope and definitions

The pilot covered September 1-30, 2026. Project owner: Maya Chen. Currency: USD. The workload contained 1,000 requests: 600 classification, 250 extraction and 150 summarization requests.

A request succeeds when the service returns a response. A response passes quality only when it meets the task-specific answer rubric. Success and quality are separate outcomes.

2. Observed pilot results

Metric	Baseline	Candidate
Requests attempted	1,000	1,000
Successful responses	990	980
Quality-passing responses	950	940
Median latency	820 ms	640 ms
95th-percentile latency	1,700 ms	1,900 ms
Total API cost, including failed attempts	$12.50	$8.00

Quality-passing responses are a subset of successful responses. In this report, quality-pass rate uses all attempted requests as its denominator. Cost per quality-passing response uses total API cost divided by quality-passing responses.

The two configurations were both executed during evaluation. Their combined cost is evaluation expenditure. The difference between their costs is a modeled saving for serving the same volume using only one configuration.

Usage and calculation details

3. Single-request billing example

Request NS-041 used the fictional model Cedar-S. Provider-reported input tokens: 12,000, including 3,000 cached-read tokens. Output tokens: 2,000, including 500 reasoning tokens. Total tokens: 14,000. Reasoning tokens are already included in output tokens.

Disjoint billable category	Tokens	USD per million
Uncached input	9,000	$0.20
Cached-read input	3,000	$0.05
Output (includes reasoning)	2,000	$0.80

No cache writes, tool charges or other charges applied to NS-041. Calculate each category as tokens multiplied by its rate divided by 1,000,000. Do not add reasoning tokens to output tokens again.

4. Revisions and missing information

Revision 1 listed candidate cost as $7.60. That figure omitted $0.40 of failed-attempt charges. Revision 2 supersedes it: the authoritative candidate total is $8.00, including all attempts.

The report does not provide measured electricity use, GPU type, provider region, October traffic, or the exact calendar date of the next pilot. No claim about those facts can be supported by this document.

5. Policy and projection assumptions

The release policy requires an all-attempt quality-pass rate of at least 94.5%. There is no rounding before applying this threshold. A lower cost alone does not satisfy the policy.

For a hypothetical month with 100,000 requests, assume exactly the same task mix, cost per attempted request and performance rates as this pilot. Assume no additional router, verification or fallback charges. This is a linear projection, not observed production spend.

Lower median latency does not imply better tail latency. Lower API cost does not prove lower physical energy use.

Questions for the app

Answer using only this PDF. Cite the source page and section for each answer. For calculations, show the formula and units. If information is missing, say "Not provided in the document." Do not use external facts.

Q01. Who owns the pilot, and what dates did it cover?

Q02. What percentage of requests were extraction tasks?

Q03. What were the candidate request-success rate and all-attempt quality-pass rate?

Q04. Does the candidate meet the release quality threshold?

Q05. What were the modeled cost saving per 1,000 requests and percentage reduction?

Q06. How much did running both configurations actually cost?

Q07. What is the candidate cost per quality-passing response, rounded to six decimal places?

Q08. Did candidate median latency and 95th-percentile latency both improve?

Q09. For NS-041, report uncached input, cached input, output, reasoning and total tokens.

Q10. Calculate the full NS-041 API cost and each component.

Q11. Should the candidate pilot cost be reported as $7.60 or $8.00, and why?

Q12. Under the stated 100,000-request projection, what are baseline cost, candidate cost and savings?

Q13. How many kWh of measured electricity did the candidate save?

Q14. What GPU type and provider region were used?

Q15. Summarize whether the candidate should be released under the stated policy in two sentences.